

*"News, rumors, revelations,
and prophecies from the age
of artificial intelligence."*

The Ai Times

Editorial Board:
Alshival.Ai



Samuel Cavazos

Top 100 Data & Analytics Professional

OpenAi spices things up with Jalapeño

OpenAi and Broadcom have partnered up to unveil a new chip optimized for LLM-inference. **The chip, named Jalapeño**, is designed to accelerate the performance of large language models, making them faster and more efficient.

The Ai Times asked ourselves: "Will this chip help improve our Kaggle leaderboard rankings?" and we have determined that it will not.

The important word is inference. Training is where a model learns; inference is where it answers. Every ChatGPT reply, Codex edit, image prompt, and customer-support bot response has to run somewhere, and today much of that work still depends on expensive, scarce Nvidia hardware. **Axios reports** that Jalapeño is aimed at customer-serving workloads rather than giant training runs. If OpenAi can make inference cheaper per response, it gets better margins, more control over its infrastructure, and a little less dependence on the global GPU line.

We are excited to see how this collaboration will impact the AI landscape and look forward to testing the chip's capabilities in real-world applications.

Tesla Sued for Fatal FSD Crash in Katy, Texas

The Ai Times recently reported on a fatal crash involving a Tesla Model 3 in Katy, Texas, where the vehicle's Full Self-Driving (FSD) system was allegedly engaged at the time of the accident. The crash resulted in the death of 76-year-old Martha Avila, who was inside her home when the car crashed into it.

The family of Martha Avila has sued Tesla and driver Michael Butler after a Tesla Model 3 crashed into a home in Katy, Texas, killing the 76-year-old grandmother inside. The crash happened on June 19. According to reporting based on the Harris County Sheriff's Office statement, Butler told investigators that an automated driving-

assistance system was engaged when the car left the roadway, entered the brick residence at high speed, and struck Avila. **NHTSA opened a special crash investigation**, and **Business Insider reported** that the exact Tesla feature in use has not yet been publicly established.

Tesla's response was unusually direct. Ashok Elluswamy, Tesla's head of AI, wrote that the driver manually overrode self-driving by pressing the accelerator to 100 percent, that the vehicle reached 73 mph, and that the accelerator remained pressed after impact. Data is not always an accurate representation of reality, as technical issues can cause discrepancies between what the vehicle logs show and what actually happened.

Elon Musk also argued that Full Self-Driving does not normally drive at full speed through residential streets. But speed control is already one of the live regulatory questions around FSD. **Reuters reporting summarized by Omni** says Sweden's Transport Administration warned the EU vehicle committee that Tesla's "speed offset" feature can let FSD drive above posted speed limits. The agency recommended voting against a broader EU rollout unless Tesla removes the feature, arguing that it could undermine both traffic law and the safety benefits automated driving is supposed to provide. **Omni also reported** that Finland and Norway share the concern, with the issue headed for discussion at the EU vehicle committee on June 30.

The big question right now is whether the vehicle Michael Butler was driving was equipped with Hardware 3 or Hardware 4. Tesla's FSD computer transition occurred in 2019, and the company has been selling cars with Hardware 4 since 2024. The difference is significant: Hardware 4 has more than double the memory bandwidth of Hardware 3, which is critical for running the neural networks that power FSD. In April, during Tesla's **Q1 2026 earnings call**, Musk acknowledged that Hardware 3 is not capable of unsupervised Full Self-Driving. **The Verge quoted him** saying that HW3 "simply does not have the capability" for unsupervised FSD, citing its much lower memory bandwidth com-

pared with Hardware 4. Tesla's own Q1 shareholder deck still described FSD as supervised, with active driver supervision required and the vehicle not autonomous. But that disclaimer does not erase the deeper problem: Tesla sold the public a feature called Full Self-Driving for years while the company's CEO now says millions of older cars cannot deliver the unsupervised autonomy customers were led to expect.

The Katy vehicle's model year, VIN, hardware generation, and software version have not been publicly identified. That must be said clearly. If the car was a refreshed 2024-or-newer Model 3, it likely used Hardware 4. If it was an older Model 3 built after Tesla's 2019 FSD-computer transition, it may have used Hardware 3. Without the VIN or Tesla's logs, nobody outside the investigation should claim to know.

But the recall question remains fair. NHTSA recalls are supposed to address safety defects or failures to meet safety standards. Tesla has already recalled FSD Beta for unsafe behavior around intersections and speed limits, and recalled Autosteer because its controls could be insufficient to prevent driver misuse. If HW3 cannot achieve the thing Tesla sold as Full Self-Driving, why is that only a customer-service problem and not a safety problem? If the answer is that FSD is merely supervised driver assistance, then the name is doing dangerous work. If the answer is that FSD is on a path to autonomy, then the public deserves to know which cars are physically capable of getting there, which are not, and why the incapable ones are still carrying the brand.

Announcements

The Ai Times is now on Substack. Scan the code below to get each new issue, along with the sources, side quests, and stray prophecies that did not fit on the page.

THE
AI
TIMES

