

"News, rumors, revelations,
and prophecies from the age
of artificial intelligence."

The Ai Times

Editorial Board:
Alshival.Ai



Samuel Cavazos

Top 100 Data & Analytics Professional

Anthropic Gets Its Hall Pass Back

Anthropic's Fable 5 is back in global circulation, which is by far one of the strangest product-launch approval process in Silicon Valley history. **Ars Technica reports** that the U.S. has lifted export curbs on Claude Fable 5 and Claude Mythos 5 about three weeks after the Trump administration flagged the models as national security risks. In those three weeks, far more powerful and potentially more dangerous models have been released.

Anthropic says the June 12 order required it to restrict access to foreign nationals, whether inside or outside the United States. As of July 1, Fable 5 is available globally across Claude Platform, Claude.ai, Claude Code, and Claude Cowork. Mythos 5 is different. Anthropic says access has been restored only to a set of U.S. organizations, with work underway to expand access to domestic and international partners in Project Glasswing, the defensive-cybersecurity program for trusted researchers. What interests me is that these models are built on principles of mathematics and data science. Anyone with enough compute power can create these models. Should we confiscate everyone's GPU and make the math forbidden knowledge? That should do it.

Indeed, we should just erase *gradient descent* from every textbook. No more talks of *transformers* and *embeddings* and *attention mechanism*.

Should we make *convolution* a sin?

Meta Is Doing Cloud Computing

During the Covid era, the people could not meet up with friends and family. They could not meet publicly. Schools were closed. Libraries were shut down. Suddenly, reality forced the people to live their lives in isolation.

But there was a silver lining. The people found a way to congregate anyway. *Facebook Live* streams started springing up everywhere, letting friends, families, churches, small businesses, and whole communities gather in the only public square they could reach. The best part

was that *Facebook Live* was free.

Then Mark gave us the *Llama* series. For free.

What... a... bro.

I do not talk about Mark Zuckerberg very much. That is because I have nothing but admiration for the things he has done. He continues to prove that despite his massive wealth, he has self-control. He isn't going to charge us for every little thing. That is a rare quality in this industry.

So how will Meta make money then?

Bloomberg now reports that Meta is building a cloud business to sell access to its AI models and excess compute. The plan is not official. Meta declined to comment, and these may be just rumors at this point, but a lot of people who seem to know what they are talking about are talking about it and so now I am talking about it too. The idea is to sell API access to Meta-hosted models, rent raw compute capacity, or do some version of both.

We are rooting for you, Mark.

Sam Altman's Essay

Sam Altman has published a **Financial Times essay**. The OpenAI chief executive argues that frontier AI needs a global framework for things like safety testing, shared standards, and dependable access for countries and companies that follow the rules.

The concrete proposal is a U.S.-led international forum, a standing body that would set standards, evaluate model capabilities and risks, certify participants, and decide who gets frontier access without launch-day chaos.

In 2023, Altman, Greg Brockman, and Ilya Sutskever wrote about the **governance of superintelligence**, arguing that systems above a certain capability threshold may eventually need a governing body. The idea was to put special controls around systems powerful enough to create national or global risks.

The Fable 5 episode at Anthropic would not have occurred with that machinery in place. At least in theory. But this connects back to that problem: you can

gate Fable 5, but you cannot gate *gradient descent*. The public is getting better at training its own models, as Palantir's **Alex Karp** recently discussed with CNBC. They may not be as capable as frontier systems, but smaller, specialized models can still be dangerous. A facial-recognition model trained to detect a particular person's face in a crowd does not need a billion-dollar cluster. It can be trained on a regular GPU and weaponized for surveillance, stalking, or combined with other technologies for automated warfare.

Altman says citizens and elected governments should make the rules, not a handful of companies in San Francisco. Correct. Also awkward. The companies in San Francisco are writing the essays, advising the government, attending the G7 meetings, building the models, and measuring the risks because they have deemed themselves responsible enough to do so. Altman's answer could become a licensing club where the largest labs, richest governments, and best-connected companies decide who is responsible enough to touch the future.

I have been working in artificial intelligence since 2016. There were not many opportunities in the field back then and if you mentioned machine-learning to someone they would look at you like you had grown a second head. It was not until Sam Altman & OpenAI made their GPT models available to the public that the field began to bloom. I know that Altman approaches with good intentions, but now he is a part of the machinery that requires governing, so maybe we should not take his essay into consideration until a governing body tells us it is safe to do so.

Announcements

The Ai Times is now on Substack. Scan the code below to get each new issue, along with the sources, side quests, and stray prophecies that did not fit on the page.

THE
AI
TIMES

